

```
import torch
import torch.nn as nn

import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import numpy as np
import gym
from collections import namedtuple
import random
import matplotlib.pyplot as plt

!pip install selenium
!apt-get update # to update ubuntu to correctly run apt i
!apt install chromium-chromedriver
!cp /usr/lib/chromium-browser/chromedriver /usr/bin
```

```
pip install gym-chrome-dino
```

```
device = torch.device('cuda' if torch.cuda.is_available())
```

Abstract

I wrote this double deep q algorithm to train an RL agent to play the chrome dino game. I was initially using the normal DQN algorithm , but the training was very unstable , SO later I switched to Double Deep q networks. The entire pipeline is I have an image encoder which is a basic cnn, where I bundle 4 frames 1 current and 3 previous and encode it using CNN , after encoding the vector will be fed into the q network as a state , and the network assigns q values, for each state action pairs. The algorithms uses replay buffer to store the experience/episodes , state, action , reward and next state values. I have then random sampled experiences and trained the network again so that the agent doesn't forget past experiences, there's also priority replay buffers where episodes which have high q values - reward differences will be sampled first. As I am using double deep q networks, I am using 2 networks policy and a target network.

I am also using epsilon - greedy strategy which is mostly common for RL algorithms which pushed the agents to explore more increasing the randomness in taking actions, once it's trained on all possibilities, we make the model exploit more with occasional exploration.

Training :

For each experience in the batch, the target Q-value is computed using the Double Q-Learning approach:

The policy network is used to determine the action taking the argmax Q-value for the next state.

The target network then evaluates the Q-value for that action to compute the target.

The Q-Network is updated by minimizing the difference between the predicted Q-values and the computed target Q-values.

The target network updates periodically using soft update, this is the main advantage which keeps the target stable.

This is the demo that i got: <https://www.youtube.com/watch?v=3nfH1JJUL0g>

```
#Install the chrome dino packages for running the gym simu
import gym_chrome_dino
from gym_chrome_dino.utils.wrappers import make_dino
import cv2
```

```
num_epochs=4
alpha=1e-4
batch_size=512
```

```
#Creating a named tuple for storing the experience
```

```
Experience = namedtuple('Experience',('state','action','n
```

```
#This is the class that I am using for replaymemory, whic
#batches to train later on, This is fine as the q learnin
#are enough experiences.
```

```
class replaymemory():
    def __init__(self,capacity):
        self.capacity=capacity
```

```
self.memory=[]
self.push_count=0

def push(self,experience):
    if(self.push_count>self.capacity):
        self.memory[self.push_count%self.capacity]=experien
    else:
        self.memory.append(experience)
    self.push_count+=1

def sample(self,batch_size):
    e=random.sample(self.memory,batch_size)
    #print(type(e1.state))
    states=torch.tensor([e1.state for e1 in e]).to(device)
    actions=torch.tensor([e1.action for e1 in e]).to(devi
    next=torch.tensor([e1.next_state for e1 in e]).to(dev
    rewards = torch.tensor([e1.reward for e1 in e]).to(
    dones = torch.tensor([e1.done for e1 in e])

    #states=states.view(64,1,80,160)
    return states.float(),actions,next,rewards,dones

def can_provide(self,batch_size):
    return len(self.memory)>=batch_size

#I am just using this to create a 4 framed array, so that

class cstate():
    def __init__(self,size):
        self.size=size
        self.states=[]
        self.push_count=0
    def push(self,state):
        if self.push_count<self.size:
            self.states.append(state)
            self.push_count=self.push_count+1
        else:
```

```
self.states.pop(0)
self.states.append(state)
```

```
#This is the class for epsilon strategy , where with gree
# with the highest q value, where as for get_action depen
# in the inital phases it favours more of exploration and
```

```
class eps_strat():
    def __init__(self,num_actions,start,decay,end,device):
        self.cur_step=0
        self.num=num_actions
        self.device=device
        self.start=start
        self.decay=decay
        self.end=end
    def get_action(self,policy,state,step):
        self.cur_step=step
        ep=self.end+(self.start-self.end) * np.exp(-1*self.cur_step)
        #self.cur_step=self.cur_step+1
        #print(ep,self.cur_step)
        if ep>random.random():
            #print("NO")
            action=random.randrange(self.num)
            return action
        else:
            policy.eval()
            with torch.no_grad():
                #print("yes")
                return torch.argmax(policy.forward(state)).item()
    def get_value(self,step):
        self.cur_step=step
        ep=self.end+(self.start-self.end) * np.exp(-1*self.cur_step)
        return ep
    def greedy(self,policy,state):
        #print("greedy")
```

```
policy.eval()  
with torch.no_grad():  
    return torch.argmax(policy.forward(state)).item()
```

#This is my convolutional network architecture

```
class convnet(nn.Module):  
    def __init__(self):  
        super(convnet,self).__init__()  
        self.conv1= nn.Conv2d(4,32,8)  
        self.conv2 = nn.Conv2d(32,64,4)  
        self.conv3=nn.Conv2d(64,64,3)  
        self.fc1=nn.Linear(64*68*68,512)  
        self.fc2=nn.Linear(512,64)  
        self.fc3=nn.Linear(64,2)  
  
    def forward(self,x):  
        out=F.relu(self.conv1(x))  
        out=F.relu(self.conv2(out))  
        out=F.relu(self.conv3(out))  
        out = out.view(-1,64*68*68)  
        out=F.relu(self.fc1(out))  
        out=F.relu(self.fc2(out))  
        out=self.fc3(out)  
        return out
```

```
class lin(nn.Module):  
    def __init__(self):  
        super(lin,self).__init__()  
        self.fc1=nn.Linear(2,128)  
        self.fc2=nn.Linear(128,64)  
        self.fc3=nn.Linear(64,2)  
    def forward(self,x):  
        out=F.relu(self.fc1(x))  
        out=F.relu(self.fc2(out))  
        out=self.fc3(out)
```

```
return out
```

```
# This is the duelling network that I wanted to try out,  
class duelling(nn.Module):
```

```
    def __init__(self,num_actions):  
        super(duelling,self).__init__()  
        self.a = num_actions  
        self.conv1 = nn.Conv2d(4,32,8)  
        self.conv2 =nn.Conv2d(32,64,4)  
        self.conv3 = nn.Conv2d(64,64,3)  
        self.fc1_value = nn.Linear(64*68*68,512)  
        self.fc2_value = nn.Linear(512,64)  
        self.fc3_value = nn.Linear(64,1)  
        self.fc1_adv = nn.Linear(64*68*68,512)  
        self.fc2_adv = nn.Linear(512,64)  
        self.fc3_adv = nn.Linear(64,self.x)
```

```
    def forward(self,x):  
        out=F.relu(self.conv1(x))  
        out = F.relu(self.conv2(out))  
        out = F.relu(self.conv3(out))  
        out = out.view(-1,64*68*68)  
        state_value = F.relu(self.fc1_value(out))  
        state_value = F.relu(self.fc2_value(state_value))  
        state_value = self.fc3_value(state_value)  
        advantage = F.relu(self.fc1_adv(out))  
        advantage = F.relu(self.fc2_adv(advantage))  
        advantage = self.fc3_adv(advantage)
```

```
        return state_value+advantage-torch.mean(advantage)
```

```
#This is the function that I am using to update the targe
```

```
def soft_update(local,target,t):  
    for target_param,local_param in zip(target.parameters()  
        target_param.data.copy_(t*local_param.data+(1-t)*targ
```

```
#Initialising the target net and policy network and makin
```

```
target_net = convnet().to(device)
policy=convnet().to(device)
soft_update(policy,target_net,1)
```

```
num_episodes=1000000
from PIL import Image
```

```
#converting the image frames by doing pre processing.
```

```
def convert1(observation):
    cv2_imshow(np.array(observation))
    ob=torch.tensor(observation)
    ob=ob.view(4,80,160)
    ob=ob.numpy()

    f_ob=[]
    for i in range(4):
        l=ob[i]
        cv2_imshow(l)
        print("wtf")
        im = Image.fromarray(l)

        f_ob.append(np.asarray(im.resize((80,80))))

    f_ob=np.array(f_ob)

    ob=torch.tensor([f_ob]).to(torch.float)
    ob=ob.view(-1,4,80,80)
    return ob
```

```
#initialising memory replay with 10000 capacity
memory = ReplayMemory(10000)
```

```
mem=replaymemory(10000)

from PIL import Image
def convert(observation):
    ob=np.array(observation,dtype="float32")
    #print(ob.shape)

    ob=cv2.cvtColor(ob,cv2.COLOR_BGR2GRAY)
    im = Image.fromarray(ob)
    ob = np.asarray(im.resize((80,80)))

    return ob

#Creating and intialising the gym environment and pushing
env = gym.make('ChromeDino-v0')
from google.colab.patches import cv2_imshow

ob=env.reset()
done_1=False
ob=convert(ob)

print(ob.shape)
stack=cstate(4)

stack.push(ob)
stack.push(ob)
stack.push(ob)
stack.push(ob)

a2=torch.tensor([stack.states])

a= nn.Conv2d(4,32,8)
b = nn.Conv2d(32,64,4)
c=nn.Conv2d(64,64,3)
out=a(a2)
out=b(torch.tensor(out))
out=c(torch.tensor(out))
```



```
print(out.shape)
#env=env.unwrapped
state=env.reset()
reward_sum=0
gamma=torch.tensor(0.999)

start=1
end=0.00001
decay=0.1
#decay=0.1

#tau for soft update

tau=0.89

t_decay=eps_strat(2,0.9,0.01,0.001,device)
print(tau)

(80, 80)
torch.Size([1, 64, 68, 68])
0.89
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:24: UserWarning:
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:25: UserWarning:

FILE='/content/drive/MyDrive/chrome_dino/target_net_dino_
File='/content/drive/MyDrive/chrome_dino/model_dino_3.pth
from IPython.display import clear_output
#model.load_state_dict(torch.load(file_path))
target_net.load_state_dict(torch.load(FILE))
policy.load_state_dict(torch.load(File))

<All keys matched successfully>
```

Double-click (or enter) to edit

```
import cv2
height=720
width=1200
fps=10
fourcc = cv2.VideoWriter_fourcc(*'mp4v') # setting up vid
```

```
FILE='/content/drive/MyDrive/chrome_dino/target_net_dino_5
File='/content/drive/MyDrive/chrome_dino/model_dino_5.pth'
step=0
step_1=0
eps=0
lo=[] # list to store loss values for plotting
re=[] # list to store rewards for plotting
video=[]
video_ep=[] # List to store video frames for the current
max_rew=0 # keeping track track of the max reward

for eps in range(num_episodes): # Training loop over a gi

    reward_sum=0 # initi total reward for the current epi
    state=env.reset() # Reseteting the environment
    state=convert(state)

    done_1=False # flag to check if the current episode i
    video_ep=[]
    while not done_1: # Loop until the episode ends
        step=step+1
        ep=eps_strat(2, start, decay, end, device) # Crea

        # Get the current state stack and reshape it for t
        a2=torch.tensor(stack.states)
        a2=a2.view(1,4,80,80)
        action= ep.get_action(policy, a2.float()).to(device

        next_state, reward, done_1, info = env.step(action
        video_ep.append(next_state) # Add frame to curren

        next_state=convert(next_state)
        stack.push(next_state) # Push the new state onto
        a3=torch.tensor(stack.states)
        a3=a3.view(1,4,80,80)

        reward = 1
```

```
if done_1:
    reward = -10 # Setting reward to -10 for game

# storing the experience in the replay memory
e1 = Experience(a2.numpy(), action, a3.numpy(), re
mem.push(e1)

reward_sum += reward # accumulate the rewards for
state = next_state #update to nest state

if max_rew < reward_sum:
    max_rew = reward_sum
    video = video_ep.copy()

re.append(reward_sum) # Track rewards for plotting
print(eps, reward_sum)

if mem.can_provide(64): # If enough experiences are i
    epochs = 10
    for i in range(epochs):
        states_1, actions, next_states, rewards, done

        # Reshape the sampled states for the network
        states_1 = states_1.view(64, 4, 80, 80)
        next_states = next_states.view(64, 4, 80, 80)

        # Calculate the expected Q values using the po
        ex = policy.forward(next_states.float()).detac
        target_q = target_net.forward(next_states.floa
        j = [t.index(max(t)) for t in ex] # Get the a

        # Select the target Q-values corresponding to
        target_q = torch.tensor([target_q[i][j[i]] for

        done = done.int().to(device)
        q = (rewards + gamma * (target_q * (1 - done))
        q = q.unsqueeze(1)
        l = torch.tensor([actions[i] for i in range(64
```

```

l = l.unsqueeze(1)

# Calculate the loss between the predicted Q-v
exp_q = policy.forward(states_1.float()).gather(1, l)
loss = nn.MSELoss()
tr = loss(exp_q, q)
lo.append(tr) # Track loss for plotting

# Perform a gradient descent step
optimizer = torch.optim.Adam(policy.parameters)
optimizer.zero_grad()
tr.backward()
optimizer.step()



if eps % 100 == 0: # Every 100 episodes, update the target net
    soft_update(policy, target_net, 1)
    episodes = range(len(lo))
    epi = range(len(re))
    plt.plot(episodes, lo, 'g', label='loss')
    plt.xlabel('episodes')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()
    plt.plot(epi, re, 'g', label='rewards')
    plt.xlabel('episodes')
    plt.ylabel('rewards')
    plt.legend()
    plt.show()

if eps % 50 == 0: # Save the model weights and best v
    torch.save(target_net.state_dict(), FILE)
    torch.save(policy.state_dict(), File)
    print("Weights saved")
    np.save('dino', np.array(video))

```

0 27





Weights saved

1 3115

2 319

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-115-ae600cf9c877> in <module>()
    37     #print(action)
    38
--> 39     next_state, reward, done_1, info = env.step(action)
    40     #cv2_imshow(next_state)
    41     #time.sleep(0.5)
```

```
----- 5 frames -----
/usr/local/lib/python3.7/dist-packages/PIL/PngImagePlugin.py in
load_end(self)
    890         logger.debug("%r %s %s (unknown)", cid, pos, length)
    891         ImageFile._safe_read(self.fp, length)
--> 892         self._text = self.png.im_text
    893         if not self.is_animated:
    894             self.png.close()
```

KeyboardInterrupt:

```
for i in video:
    cv2_imshow(i)
    time.sleep(0.07)
```

```
time.sleep(0.01)  
clear_output()
```